

Solution-Operator Semi-Discretization: General-Order Multiplication-Free Stability Analysis of Time-Periodic Delayed Systems and the Accuracy–Sparsity Trade-Off

Journal of Vibration and Control
XX(X):1–12
©The Author(s) 2026
Reprints and permission:
sagepub.co.uk/journalsPermissions.nav
DOI: 10.1177/ToBeAssigned
www.sagepub.com/

SAGE

Daniel Bachrathy¹

Abstract

The multiplication-free semi-discretization method (MFSD) reduces the stability analysis of time-periodic delay differential equations to a sparse, banded generalized eigenvalue problem with linear time complexity in the resolution p , but its piecewise-constant step limits the convergence order to two. This paper introduces the *solution-operator semi-discretization* (SOSD) method: the integrator inside each step becomes a free design choice, and the residual block rows of any explicit Runge–Kutta or implicit collocation-type scheme (Gauss–Legendre, Radau, Lobatto) embed in the banded system through a Kronecker-product structure of the Butcher tableau and the system matrices. The full superconvergent order $2s$ of s -stage Gauss collocation is observed for mesh-commensurate delays, and orders between $s+1$ and $2s$ are observed for general (including time-periodic) delays. Because the cost at fixed step count grows polynomially with the stage number while h -refinement stays essentially linear, the CPU-time optimum is a problem-dependent *sweet spot*, located here for systems requiring tens to many hundreds of samples per period. A Shannon-type sampling bound turns the measurements into an a priori recipe for choosing the stage and step counts, demonstrated end-to-end on high-resolution stability charts. All methods are implemented in the open-source Julia package SOSD.jl.

Keywords

time-delay systems, semi-discretization, solution operator, Runge–Kutta methods, Floquet theory, stability analysis, computational complexity

1 Introduction

Time-periodic linear delay differential equations (DDEs) of the form

$$\dot{\mathbf{x}}(t) = \mathbf{A}(t) \mathbf{x}(t) + \sum_{k=1}^g \mathbf{B}_k(t) \mathbf{x}(t - \tau_k(t)) + \mathbf{c}(t), \quad (1)$$

with T -periodic coefficients $\mathbf{A}(t)$, $\mathbf{B}_k(t)$, delays $\tau_k(t)$ and forcing $\mathbf{c}(t)$, govern a wide range of engineering and biological systems: regenerative machine-tool chatter (Tobias 1965; Altintas 2012; Munoa et al. 2016), turning with spindle-speed variation (Long and Balachandran 2007; Seguy et al. 2008), feedback control with latency (Stépán 1989), and seasonally forced population dynamics. Their asymptotic stability is determined by the dominant characteristic multipliers of the monodromy operator associated with one period (Insperger and Stépán 2011).

Two families of numerical methods dominate practice. Semi-discretization (SD) (Insperger and Stépán 2002, 2011) advances the state with a fixed step $\Delta t = T/p$, treating the delayed term (and, in the classical zeroth- and first-order variants, the coefficients) as piecewise constant or piecewise linear per step; its convergence order is limited to two (Insperger et al. 2008). Collocation and pseudospectral methods (Breda et al. 2005, 2015) represent the solution segment by a global high-order polynomial

and converge spectrally for smooth problems. A recurring claim in the literature is that collocation “outperforms” semi-discretization; we argue in Section 4 that such comparisons are structurally unfair, because they compare an order-refinement strategy (p -type refinement in FEM terminology, raising the polynomial order) with an h -refinement strategy (shrinking the step) — the same category error as declaring p -type finite elements universally superior to h -type ones.

The recent multiplication-free semi-discretization method (MFSD) (Bachrathy 2026) removed the main cost bottleneck of SD: instead of multiplying p one-step matrices into a dense monodromy matrix ($\mathcal{O}(p^{2.8})$ in practice), every step is written as a residual equation and all p steps are stacked into one sparse banded system, split into the one-period mapping $\Phi_{LY_p} = \Phi_{RY_0}$; with the Floquet ansatz $\mathbf{y}_p =$

¹Department of Applied Mechanics, Faculty of Mechanical Engineering, Budapest University of Technology and Economics, Műegyetem rkp. 3., H-1111 Budapest, Hungary.

Corresponding author:

Daniel Bachrathy, Department of Applied Mechanics, Faculty of Mechanical Engineering, Budapest University of Technology and Economics, Budapest, Hungary.

Email: bachrathy@mm.bme.hu

$\mu \mathbf{y}_0$ the multipliers follow from the generalized eigenvalue problem $\mu \Phi_L \mathbf{v} = \Phi_R \mathbf{v}$ at $\mathcal{O}(p)$ cost.

The present paper starts from a simple observation: *the multiplication-free form only needs a per-step residual*. Semi-discretization is, structurally, a fixed-step time integrator whose only approximation is the per-step treatment of the coefficients and the delayed term. Consequently, the piecewise-constant step can be replaced by *any* one-step scheme — an explicit Runge–Kutta method or an implicit collocation-type method — by placing the Butcher-tableau coefficients into the residual block rows and augmenting the per-step state with the internal stage values. This raises the achievable convergence order arbitrarily, at a quantifiable structural price: the sparse matrices become denser as the stage number grows. We call the resulting family the *solution-operator semi-discretization* (SOSD): the banded pair (Φ_L, Φ_R) is a sparse representation of the one-period solution operator of (1), and the per-step integrator merely selects the order of that representation.

The contributions, precisely delimited, are the following:

1. the systematic embedding of arbitrary Runge–Kutta and collocation steps (via their Butcher tableaux and continuous extensions) into the multiplication-free banded stability framework for time-periodic DDEs, including time-periodic delays (Section 3);
2. the explicit-versus-implicit matrix-structure and CPU trade-off analysis, and the demonstration that the single-step limit reproduces the known global collocation method (Section 3.3);
3. the identification and characterization of the accuracy–sparsity *sweet spot*: for a given accuracy target the CPU optimum lies at moderate per-step order and many steps, at a location that depends on the problem dimension and resolution demand (Section 5);
4. a fair-comparison methodology that separates order refinement from h -refinement and reports work–precision curves with variance on one fixed machine configuration (Section 4); and
5. an open-source Julia implementation, SOSD.jl, with all scripts needed to reproduce every figure.

Neither collocation itself (Breda et al. 2005) nor the multiplication-free assembly (Bachrathy 2026) is claimed as new; the value of the paper is the unifying embedding, the trade-off quantification, and the practical guidance it yields.

2 Background: MFSD and semi-discretization as fixed-step integration

Discretize one principal period into p uniform steps, $\Delta t = T/p$, $t_n = n \Delta t$, and let $r = \lceil \tau_{\max}/\Delta t \rceil$ count the steps covering the largest delay. Classical SD constructs a one-step mapping $\mathbf{x}_{n+1} = \mathbf{P}_n \mathbf{x}_n + \mathbf{R}_n \mathbf{x}_{n-r}$ from the exact solution of the frozen-coefficient system over $[t_n, t_{n+1}]$ (Insperger and Stépán 2011). The monodromy matrix is the product of p such mappings acting on the discretized history $\mathbf{y}_n = (\mathbf{x}_n^\top, \dots, \mathbf{x}_{n-r}^\top)^\top$; forming this product scales polynomially in the resolution ($\approx \mathcal{O}(p^{2.8})$ measured in Bachrathy 2026), and its avoidance is the point of MFSD: each step contributes the residual row $\mathbf{0} = \mathbf{I} \mathbf{x}_{n+1} - \mathbf{P}_n \mathbf{x}_n - \mathbf{R}_n \mathbf{x}_{n-r}$ to a sparse

banded block system over the extended vector spanning $[t - \tau_{\max}, t + T]$, split into left/right blocks so that, with the Floquet ansatz $\mathbf{y}_p = \mu \mathbf{y}_0$, the spectrum follows from $\mu \Phi_L \mathbf{v} = \Phi_R \mathbf{v}$ via sparse Krylov iteration (Haegeman 2021) at $\mathcal{O}(p)$ total cost (Bachrathy 2026).

Two properties of this construction matter for what follows. First, the one-step map is a *fixed-step integrator*: the only approximation is that $\mathbf{A}(t)$, $\mathbf{B}(t)$ and the delayed argument are frozen (or low-order interpolated) across Δt . It is exactly this per-step treatment that caps the convergence order at two: as shown by Insperger et al. (2008), a first-order approximation of the delayed term is already optimal for the piecewise-constant coefficient scheme, and no refinement of the classical SD step exceeds order two. Second, the assembly touches the integrator only through the per-step residual; nothing in the banded structure, the left/right splitting for $p - 1 \geq r$, or the Krylov solver depends on *which* integrator produced the residual. The integrator is a free design choice — the theme of this paper.

3 Higher-order multiplication-free stepping

3.1 Stage-augmented residual blocks from a Butcher tableau

Consider an s -stage Runge–Kutta method with tableau $(\mathbf{a}, \mathbf{b}, \mathbf{c})$, a_{ij} , b_j , c_j , applied to (1) on $[t_n, t_n + \Delta t]$. Writing the stage *values* $\mathbf{Y}_i \approx \mathbf{x}(t_n + c_i \Delta t)$ (rather than stage slopes), the step reads

$$\mathbf{Y}_i = \mathbf{x}_n + \Delta t \sum_j a_{ij} [\mathbf{A}(t_{nj}) \mathbf{Y}_j + \mathbf{d}_j], \quad i = 1, \dots, s, \quad (2)$$

$$\mathbf{x}_{n+1} = \mathbf{x}_n + \Delta t \sum_j b_j [\mathbf{A}(t_{nj}) \mathbf{Y}_j + \mathbf{d}_j], \quad (3)$$

with $t_{nj} = t_n + c_j \Delta t$ and the delayed-plus-forcing terms

$$\mathbf{d}_j = \sum_{k=1}^g \mathbf{B}_k(t_{nj}) \tilde{\mathbf{x}}(t_{nj} - \tau_k(t_{nj})) + \mathbf{c}(t_{nj}), \quad (4)$$

where $\tilde{\mathbf{x}}$ denotes the interpolated delayed state (Section 3.2). The per-step unknown block is the stage-augmented vector

$$\mathbf{w}_n = (\mathbf{Y}_1^\top, \dots, \mathbf{Y}_s^\top, \mathbf{x}_{n+1}^\top)^\top \in \mathbb{R}^{(s+1)d}, \quad (5)$$

and the global unknown $\mathbf{q} = (\mathbf{w}_1^\top, \dots, \mathbf{w}_p^\top)^\top$ collects all steps of one period. Equations (2)–(3) are linear in the current and delayed states, so each step contributes one block row to a sparse system $\Phi_L \mathbf{q} = \Phi_R \mathbf{q}_{\text{hist}}$, exactly as in MFSD but with $(s+1)d$ -sized blocks.

The structure of these blocks is a Kronecker (dyadic) composition of three independent ingredients: the Butcher tableau, the system matrices, and the interpolation weights. Collecting the tableau into the compound $(s+1) \times s$ matrix

$$\hat{\mathbf{A}} = \begin{bmatrix} \mathbf{a} \\ \mathbf{b}^\top \end{bmatrix}, \quad \mathbf{K} = \hat{\mathbf{A}} \otimes \mathbf{I}_d, \quad (6)$$

and writing $\mathcal{A}_n = \text{diag}(\mathbf{A}(t_{n1}), \dots, \mathbf{A}(t_{ns}))$ for the stage-wise coefficient block, the step (2)–(3) reads, for the stacked unknown block $\mathbf{w}_n$,

$$[\mathbf{I} - \Delta t \mathbf{K} \mathcal{A}_n \mathbf{S}] \mathbf{w}_n = (\mathbf{1}_{s+1} \otimes \mathbf{I}_d) \mathbf{x}_n + \Delta t \mathbf{K} \mathbf{d}_n, \quad (7)$$

where $\mathbf{S} = [\mathbf{I}_{sd} \ \mathbf{0}]$ selects the stage part of $\mathbf{w}_n$ and $\mathbf{d}_n$ stacks the terms (4). Each delayed sample in $\mathbf{d}_n$ is itself a dyad: stage j of delay k contributes

$$\Delta t (\hat{\mathbf{A}} \mathbf{e}_j \otimes \mathbf{B}_k(t_{nj})) (\ell(\theta_{nj})^\top \otimes \mathbf{I}_d), \quad (8)$$

i.e. the outer product of a Butcher-tableau column, the delay coefficient matrix, and the continuous-extension weight row $\ell(\theta)^\top$ evaluated at the fractional position θ_{nj} of the delayed instant. The entire assembly therefore separates into precomputable factors (tableau, weights) and pointwise coefficient evaluations $\mathbf{A}(t^*)$, $\mathbf{B}_k(t^*)$ — no matrix products between steps ever occur. Figure 1 shows the resulting sparsity patterns: Φ_L has identity diagonal blocks and is block-lower triangular for *both* explicit and implicit tableaux (after the per-step elimination below), so linear solves reduce to forward substitution with zero fill-in; all delay couplings that reach into the pre-period history populate Φ_R . The per-block fill of the delay band grows with the stage number — the mechanism that will drive the accuracy–sparsity sweet spot.

In the implementation, the stage equations (2) are *pre-eliminated* per step: the stage-coupling matrix $\mathbf{M}_n = \mathbf{I} - \Delta t (\mathbf{a} \otimes \mathbf{I}_d) \text{diag}(\mathbf{A}(t_{n1}), \dots, \mathbf{A}(t_{ns}))$ of size $sd \times sd$ is factorized once (LU; a dense inverse for small sd), and the responses of $(\mathbf{x}_{n+1}, \mathbf{Y}_1, \dots, \mathbf{Y}_s)$ to (i) $\mathbf{x}_n$, (ii) each delayed-state sample, and (iii) the additive forcing are precomputed as $(s+1)d \times d$ transition blocks. This keeps Φ_L *unit block-lower triangular* — its LU factorization is trivial and fill-free — while the per-step elimination costs $\mathcal{O}(p(sd)^3)$, linear in p . For explicit tableaux (strictly lower-triangular $\mathbf{a}$) the elimination degenerates to forward substitution; for implicit collocation tableaux (Gauss–Legendre of order $2s$, Radau IIA of order $2s-1$, Lobatto IIIA of order $2s-2$ (Hairer and Wanner 1996; Butcher 2016)) it is a genuine dense solve, which is precisely where the superconvergent order is purchased.

3.2 Delay interpolation by continuous extension

The delayed argument $t_{nj} - \tau_k(t_{nj})$ generally falls between grid nodes, inside some earlier step m . Interpolating the delayed state with an order lower than that of the integrator silently caps the observed convergence order — the higher-order analogue of the remark of Insperger et al. (2008) for the classical scheme, and a classical result of DDE numerics (Bellen and Zennaro 2003). The consistent choice is the *continuous extension* of the step itself: for collocation methods, the collocation polynomial through $(\mathbf{x}_m, \mathbf{Y}_{1,m}, \dots, \mathbf{Y}_{s,m}, \mathbf{x}_{m+1})$ evaluated at the fractional position $\theta \in [0, 1]$ of the delayed instant inside step m ; for explicit methods, the scheme’s dense-output polynomial (for RK4 the classical cubic Hermite-type formula). The interpolation weights depend only on θ and the tableau, so they join the precomputed transition blocks without changing the banded structure; each delayed sample couples block row n to the two neighbouring history blocks m and $m+1$. Provided every delay spans at least one step ($\tau_k \geq \Delta t$; Appendix E treats sub-step delays), the delayed instant lies in a strictly earlier step, so these couplings stay below the diagonal and Φ_L remains triangular; couplings reaching

into the pre-period history enter Φ_R instead. Time-periodic delays $\tau_k(t)$ are handled naturally, because θ and the target block m are evaluated per stage from $\tau_k(t_{nj})$.

3.3 The single-step limit is global collocation

Setting $p = 1$ with an s -stage Gauss tableau makes the single step span the entire period: the stage values *are* then the collocation samples of a global polynomial, and the construction reproduces the known collocation/pseudospectral method for DDEs (Breda et al. 2005, 2015) in its strong form. This extreme is explicitly *not* new; we include it (i) to make the equivalence concrete — in Section 5 the single-step spectrum matches the converged reference to near machine precision — and (ii) to expose its cost: the single-step matrix is effectively dense, and the number of nonzeros of the assembled system grows as $\mathcal{O}(d^2 s^2)$ with the stage count (and its factorization cost as $\mathcal{O}(d^3 s^3)$), so pushing all accuracy into one step abandons exactly the sparsity that makes the multiplication-free form fast. The continuum between “ p large, s small” and “ $p = 1$, s large” is the design space this paper maps.

4 Fair-comparison methodology

Comparisons between DDE stability methods frequently pit an order-refinement method (collocation at increasing polynomial order) against an h -refinement method (semi-discretization at decreasing step size) and report accuracy at a single, often low, resolution on linear axes. Such comparisons cannot identify asymptotic rates, hide constant factors, and — like comparing p -type against h -type FEM refinement — do not license a universal winner. The predecessor paper (Bachrathy 2026) already criticized sub-100-resolution, non-log-log comparisons; we extend that critique into a positive prescription:

1. **Work–precision axes.** The primary comparison is CPU time versus eigenvalue error, both logarithmic, over wide ranges; the secondary one is CPU time versus resolution p with fitted scaling exponents.
2. **Repetition-robust timing.** Every timed point is the *minimum* of repeated warm-started runs (the standard BenchmarkTools practice: the minimum is immune to garbage-collection and scheduler interference, whose one-sided spikes contaminate means); repetition counts and scatter are recorded alongside the results.
3. **Both refinement families, both complexity laws.** Every method is shown under h -refinement (fixed order, growing p), where the cost of the banded assembly is *linear*, $t \propto p^1$, and under order refinement (fixed p , growing stage number s), where the cost is *polynomial* in the order — per-step elimination $\mathcal{O}((sd)^3)$ plus stage-quadratic fill. A comparison that shows only one direction hides one of the two laws. The complete picture is the (s, p) -map of Section 5.4: error and CPU time rendered over the whole plane, with cells beyond a CPU budget left uncomputed (NaN).
4. **Fixed environment.** One machine, one BLAS thread configuration, pinned package versions, deterministic start vectors; all reported alongside the results.




Figure 1. Sparsity patterns of the solution-operator pair (Φ_L, Φ_R) on the delayed Mathieu equation ($p = r = 8$): explicit RK4 (top) versus Gauss GL3 (bottom). Φ_L is unit block-lower-triangular in both cases (forward-substitution solvable, no fill-in); the delay band sits in Φ_R , with denser per-block coupling for the collocation-type stages.

5. **Verified references.** Error is measured against a reference that two distinct resolutions of a higher-order method agree on, and the reference method is cross-validated against an independent implementation (SemiDiscretizationMethod.jl).

5 Numerical results

5.1 Test systems

Four systems span the design space (parameters in Appendix A):

1. **Delayed Mathieu equation** ($d = 2$, $\tau = T = 2\pi$): the canonical smooth validation case (Insperger and Stépán 2011).
2. **Seasonal scalar model** ($d = 1$, $\tau = 2$, $T = 2\pi$): a smooth, cheap system where high per-step order excels.
3. **Turning with spindle-speed variation (SSV)** ($d = 2$, time-periodic delay, $T/\tau_{\max} \approx 9.1$): exercises the $p \gg r$ partitioning and time-varying delay handling.
4. **FEM beam with delayed boundary feedback** ($d = 28$, $\tau/T = 1/2$), optionally with act-and-wait switching of the feedback (Insperger 2006) (non-smooth stress test): the high-dimensional, high-resolution case where banded moderate order should win.

5.2 Order verification

Figure 2 and Table 1 confirm on the delayed Mathieu equation that each embedded *collocation* integrator attains its theoretical convergence order in the dominant multiplier

on this *mesh-commensurate* test ($\tau = T = p \Delta t$): the Gauss family exhibits the superconvergent order $2s$ (fitted slopes 1.96, 3.98, 5.98 and 9.98 for $s = 1, 2, 3, 5$; GL8 reaches the 10^{-16} relative floor from the coarsest measured resolutions), Radau IIA attains $2s - 1$ and Lobatto IIIA $2s - 2$. The first-order RK1 curve is included in the figures for completeness. Commensurability matters: the continuous extension of an s -stage collocation step is a degree- s polynomial with uniform accuracy $\mathcal{O}(\Delta t^{s+1})$, so for delays that fall *inside* steps, classical DDE interpolation theory (Bellen and Zennaro 2003) caps the attainable order at $s+1$ unless error cancellation occurs. The measurements bracket exactly this range: on the seasonal model ($\tau = 2$, never mesh-aligned) GL2 and GL3 reduce to fitted orders 2.9 and $3.8 \approx s+1$, while on the SSV turning system (time-varying delay) GL2 retains the full 4.0 and GL3 attains 5.3 — between $s+1$ and $2s$. Convergence curves for both non-commensurate cases, with the $[s+1, 2s]$ guide bands drawn explicitly, are collected in Appendix B (Figure 7). The explicit family is capped harder: the observed order is limited by its continuous extension, so RK4 with its classical third-order dense output converges with slope ≈ 3.2 instead of 4, and RK5 with a generic stage-value interpolant collapses to slope ≈ 2.3 — the delayed lookup, not the step, is the binding constraint, exactly as anticipated in Section 3.2. Even with the $s+1$ cap, collocation stages carry a consistent high-order extension at no extra cost, and their nodal order can exceed it up to $2s$; sustaining comparable orders explicitly would require bespoke dense output of matching order. This asymmetry is a structural argument for implicit collocation steps inside the multiplication-free framework that the usual “explicit steps are cheaper” intuition misses.

Table 1. Order verification on the delayed Mathieu equation: nominal versus fitted convergence order of the dominant multiplier (log–log slope over the pre-floor error range).

Method	Nom.	Fitted	Comment
SD (order 2)	2	3.1	aligned-delay cancellation
RK2 (Heun)	2	1.8	
RK4	4	3.2	capped: 3rd-order dense output
RK5	5	2.3	capped: stage interpolant
GL1 (midpoint)	2	2.0	
GL2	4	4.0	superconvergence 2s
GL3	6	6.0	
GL5	10	10.0	
GL8	16	—	at 10^{-16} floor, $p \geq 16$
Radau IIA, $s=3$	5	4.9	
Lob. IIIA, $s=3$	4	4.0	
Lob. IIIA, $s=4$	6	6.0	




Figure 2. Order verification on the delayed Mathieu equation: eigenvalue error versus resolution p (log–log), one curve per method, fitted slopes in the legend (explicit methods dashed).

5.3 Work–precision and time complexity

Figure 3 shows the three canonical panels for the delayed Mathieu equation: error versus p , CPU time versus p , and the work–precision plane. Every embedded integrator preserves the near-linear time complexity of the multiplication-free assembly — the fitted scaling exponents lie between 0.9 and 1.2 (median 1.05) for all methods from RK2 through GL8, rising to at most 1.6 on the high-dimensional beam — so raising the per-step order shifts the CPU-time line by an essentially constant factor (the block fill) without changing its slope. On the work–precision plane this is decisive: GL5 and GL8 reach relative errors of 10^{-14} within milliseconds, whereas the order-2 baseline (labelled SDM-O2 in the figures) spends an order of magnitude more time to reach only 10^{-6} . The same qualitative picture holds for the scalar seasonal model and for the spindle-speed-variation turning system (time-periodic delay, $p \gg r$ partitioning); the beam work–precision numbers appear in Section 5.6 and Appendix B.

For baseline parity with the predecessor (Bachrathy 2026), Figure 4 reproduces, in the present environment, the time-complexity contrast between the traditional monodromy construction (SD-classic, fitted exponent 2.2 over the affordable range — the series stops when a solve exceeds the 2 s budget, at $p = 10^3$ with 5.4 s) and the banded LR mapping (fitted exponent 0.98, steepening to ≈ 1.2 beyond

$p = 10^4$ as memory effects enter; continued to $p = 10^6$ in 13 s); both agree in the multiplier to 10^{-15} , and at $p = 10^3$ the measured times are 4.5 ms versus 5.4 s — a factor of 1200.

5.4 The accuracy–sparsity sweet spot

For a fixed relative accuracy target ε , sweeping the per-step order s while choosing the smallest resolution p that reaches ε produces a U-shaped CPU-time curve: low orders need excessive p , high orders pay polynomially growing per-step cost — stage elimination $\mathcal{O}((sd)^3)$ plus fill, with measured fixed- p exponents $t \propto s^k$, $k = 1.3$ (Mathieu, $d = 2$), 1.8 (beam, $d = 28$; map not shown) and 2.0 (turning) for rank-one delay coupling, and $k \rightarrow 3$ asymptotically as the stage elimination dominates. A first-order cost model, $\text{cost} \propto p s^k$ with $\varepsilon \propto (\Delta t/T_c)^{2s}$ for a characteristic time T_c (taking the best-case nodal order $2s$; reduced orders for non-commensurate delays raise the optimum appreciably — roughly doubling s^* — which explains the high-order optimum observed at the extreme accuracy target, and the factorial decay of the Gauss error constants is neglected), predicts the optimum

$$s^* \approx \frac{1}{2k} \ln(1/\varepsilon), \quad (9)$$

i.e. even a 10^{-12} accuracy target asks for only $s = 7$ –11 Gauss stages — moderate order, many steps. This asymptotic model applies when the fill actually binds. For the small, smooth Mathieu system it does not: the measured optima sit at $s^* = 5$ –8 with as few as $p^* = 2$ –6 steps (e.g. $\varepsilon = 10^{-12}$ in 4 ms with $s = 8$, $p = 6$, $\text{nnz}(\Phi_L) = 288$) — for low-demand problems the whole matrix is tiny and the near-spectral corner is optimal. The fill term takes over exactly in the contrasting cases: the measured per-step fill grows from 12 nonzeros at $s = 1$ to 112 at $s = 20$ on the Mathieu system but from $1.7 \cdot 10^3$ ($s = 1$) to $6.4 \cdot 10^3$ ($s = 6$) per step on the $d = 28$ beam, whose optima correspondingly retreat to $s^* = 4$ –5 (relative accuracy 10^{-8} in 11 ms at $s = 5$, $p = 6$; targets beyond 10^{-8} are excluded as they approach the reference accuracy).

Figure 5 renders the full picture: relative error and CPU time as colour maps over the (s, p) plane, with the two complexity laws visible as the gradient directions — along p at fixed s the cost grows linearly, along s at fixed p polynomially — and the per-target optima marked. Cells whose cost exceeds the CPU budget are not computed (blank). For a given error contour, the cheapest cell is never the single-step/ultra-high-order corner once the resolution demand or the system dimension is nontrivial: the polynomial s -law meets the linear p -law at moderate order.

5.5 Contrasting resolution demands: where the banded route wins

The location and the practical consequences of the sweet spot depend on how many points per period the problem physically requires. Two showcases bracket the range:

- **Low demand (tens of points).** For the delayed Mathieu equation and the seasonal scalar model, resolutions of $p \approx 10$ –30 already reach high accuracy with moderate Gauss orders; matrices are so small that




Figure 3. Work-precision study on the delayed Mathieu equation. Left: eigenvalue error vs. resolution p (clipped at $3 \cdot 10^0$ — pre-resolution values carry no information); middle: CPU time vs. p , all fitted exponents ≈ 1.0 ; right: the work-precision plane (down and left is better). CPU times are minima of repeated warm-started benchmark measurements; their relative scatter is below a few per cent and is therefore not drawn.




Figure 4. Baseline parity with the predecessor: CPU time of the traditional monodromy build (SD-classic, capped at 2 s per solve) versus the banded multiplication-free mapping (continued to $p = 10^6$) on the delayed Mathieu equation; minimum of repeated warm-started measurements, dashed lines: fitted scaling.

sparsity is irrelevant, and even the single-step spectral corner is viable.

- **High demand (hundreds of points and more).** The spindle-speed-variation turning system oscillates with $\omega_{\max}T/2\pi \approx 36$ (its fastest natural period is short relative to the principal period), so *any* method needs $N = p(s+1) \gtrsim 200\text{--}400$ samples per period to resolve the dynamics at all. Here the banded moderate-order route wins decisively at engineering accuracies: the measured optima are $(s^*, p^*) = (7, 47)$ for $\epsilon = 10^{-4}$ (21 ms), $(5, 159)$ for 10^{-6} (37 ms) and $(7, 100)$ for 10^{-8} (49 ms) — tens to hundreds of steps at moderate order, with the sample count pinned near the bound ($N \approx 380\text{--}950$). Only at the extreme target $\epsilon = 10^{-12}$ does the optimum shift towards high order, $(20, 22)$ at 91 ms — and even there $N = 462$ respects the bound and the step count remains far from the single-step corner, while no single-step cell reaches any target within the stage range.

The resolution demand can be estimated a priori by a sampling (Shannon-type) argument. The discretization carries $N = p(s+1)$ solution samples per period; resolving the fastest natural oscillation of the system, $\omega_{\max}$, requires

at least the Nyquist floor of two samples per oscillation period and, for engineering accuracy, empirically $c_s \approx 5\text{--}10$ samples:

$$N = p(s+1) \gtrsim c_s \frac{\omega_{\max}T}{2\pi}. \quad (10)$$

For the turning system ($\omega_{\max}T/2\pi \approx 36$) this predicts $N \gtrsim 200\text{--}400$, matching the observed onset of convergence; below it the computed multipliers are meaningless (for GL2 at $p = 50$, i.e. $N = 150$ — and a fortiori for the order-2 baseline, whose sample count is only p , they are off by orders of magnitude). The bound (10) sharpens the sweet-spot analysis: when N is pinned by the oscillation content rather than by the accuracy target, the cost of the banded route scales as $p s^k \approx N s^{k-1}$ — *still increasing in the order at fixed sample count* for every measured exponent k — so the smallest stage number that meets the accuracy target is CPU-optimal, and the single-step spectral corner (all N samples in one dense block) is maximally penalized. Conversely, when the problem is easy (N shrinks exponentially with growing s), higher order pays until the fill takes over. The same contrast governs high-dimensional systems: for the FEM beam ($d = 28$), memory — the per-step fill — caps the affordable stage number before CPU time does.

An important caveat on superconvergence follows. Figure 6 shows the order-refinement direction explicitly: the error versus the *total sample count* $N = p(s+1)$, sweeping the stage number at the fixed step counts $p = 1, 10, 20, 30, 40$, side by side. Superconvergence does *not* mean that a handful of stages always suffices: below the resolution bound (10) — the shaded band — the discrete spectrum is simply wrong (relative errors of order one), and the steep decay only begins once N crosses the band. Plotted against N , the message is direct: the resolution that matters is the sample count, not p or s alone. For the Mathieu system the band sits at $N \approx 10\text{--}20$ and every curve dives at or just beyond it; for the turning system it sits at $N \approx 200\text{--}400$, all curves stay at $\epsilon = \mathcal{O}(1)$ until they reach it, and a single-step scheme ($p = 1$) cannot cross it at all within the practical stage range. Raising the order buys nothing until the sampling limit is met; beyond the limit it buys everything. The order-verification and work-precision figures (error versus p at fixed s , Figures 2 and 3, and Appendix B




Figure 5. The (s, p) -map for the high-resolution-demand turning system ($p = 1 \dots 4 \cdot 10^3$, $s = 1 \dots 30$; beyond $s \approx 30$ the product-form evaluation of the continuous-extension weights loses floating-point accuracy — a barycentric implementation would extend the range, which is irrelevant in practice: all engineering-target optima lie at $s \leq 8$, and even the extreme 10^{-12} -target optimum sits at $s^* = 20$, well inside the range). Left: colour encodes CPU time; middle: colour encodes the eigenvalue error; right: both contour families overlaid — following, e.g., the $\varepsilon = 10^{-6}$ error contour (solid), the smallest CPU-time contour (dashed) it touches identifies the optimal setting, marked by the star. h -refinement (vertical) is linear in cost, order refinement (horizontal) polynomial; blank cells exceed the 5 s CPU budget. The corresponding maps for the low-demand Mathieu system are in Appendix B.

for the hard case) show the complementary h -refinement direction of the same surface, and the measured CPU-versus- s exponents appear in the (s, p) -map of Figure 5.

5.6 Explicit versus implicit steps

The conventional trade-off — explicit steps are cheap but order-capped, implicit stages are expensive but superconvergent — is reshaped by two findings. First, after per-step pre-elimination *both* families yield a unit block-lower-triangular Φ_L , so the downstream solve cost differs only through the block fill; the explicit advantage reduces to the cheaper (forward-substitution) per-step elimination, which is a linear-in- p term either way. Second, the delay-interpolation order requirement (Table 1) penalizes exactly the explicit family: sustaining order $q > 3$ explicitly requires bespoke dense output of order q , i.e. additional stages and fill, whereas the collocation families carry their matching-order continuous extension for free. In the measured work–precision curves the explicit methods are therefore competitive only in the low-accuracy regime ($\varepsilon \gtrsim 10^{-4}$), and Gauss stages dominate everywhere beyond it. On the FEM beam the contrast is stark: GL3 reaches $\varepsilon = 10^{-4}$ in 6 ms and 10^{-6} in 11 ms, whereas the order-2 baseline needs 140 ms and 795 ms ($23 \times 72 \times$), explicit RK5 needs 3.9 s for 10^{-6} ($\approx 360 \times$ slower than GL3), and $\varepsilon = 10^{-8}$ is reached by the Gauss variants alone — the explicit family is additionally barred from coarse grids by its Courant–Friedrichs–Lewy (CFL)-like stability bound on this stiff system (Appendix B).

5.7 Remark on scope: smooth systems

The high-order convergence results above — like those of every spectral or high-order method — presuppose (piecewise-)smooth coefficients with any discontinuities aligned to grid nodes. For genuinely non-smooth systems (act-and-wait switching, interrupted cutting) *all* methods of any nominal order fall back to low observed order when the discontinuity falls inside a step; this is verified for the act-and-wait beam in Appendix C. Two practical consequences: the sweet-spot recommendation of moderate order with many steps holds *a fortiori* (the sparsest variants lose nothing

where high order cannot pay off), and the linear-complexity advantage of the solution-operator assembly is unaffected. A discontinuity-aware treatment (mesh alignment, step splitting) is outside the scope of this paper.

6 Discussion: practical guidance

The measurements condense into a short decision rule for practitioners — an a priori recipe requiring only the fastest natural frequency $\omega_{\max}$ of the system, the period T , and the accuracy target ε :

1. **Resolution bound:** $N_{\min} \approx c_s \omega_{\max} T / (2\pi)$ with $c_s \approx 5\text{--}10$ (Eq. (10));
2. **Stage number:** $s \approx \max(2, \lceil \ln(1/\varepsilon) / (2k) \rceil)$ with the measured cost exponent $k \approx 1.3\text{--}2.0$ (Eq. (9)), capped by memory for large d ;
3. **Step count:** $p = \lceil N_{\min} / (s+1) \rceil$, then one or two h -refinements to confirm the target — cheap, since the cost is linear in p .

In more detail:

1. **Default: Gauss stages, moderate order.** The stage-number rule of step 2 above — $s \approx \lceil \ln(1/\varepsilon) / (2k) \rceil$ with the measured $k \approx 1.3\text{--}2.0$ — gives $s = 7\text{--}11$ even for $\varepsilon = 10^{-12}$, in agreement with the measured optima $s^* = 4\text{--}8$ at engineering targets across all systems. Then pick p by one or two refinement steps until the target is met — the cost is linear in p , so overshooting p is cheap, whereas overshooting s inflates the per-step cost polynomially.
2. **Large systems: cap the order earlier.** The per-step fill grows with d^2 and polynomially with s ; for the FEM-scale beam ($d = 28$) the measured optima sit at $s^* = 4\text{--}5$ with $p^* \leq 6$, and memory, not CPU, becomes the binding constraint for larger stage numbers.
3. **Avoid the spectral corner for engineering models.** The single-step (global collocation) mode is elegant for small, smooth, low-accuracy-run-count problems, but it abandons sparsity ($\text{nnz} \sim d^2 s^2$ in one dense




Figure 6. Order refinement at the fixed step counts $p = 1, 10, 20, 30, 40$: relative error versus the total sample count $N = p(s+1)$ for the easy (left) and hard (right) system, error clipped at $3 \cdot 10^0$. The shaded band is the sampling estimate (10) with $c_s = 5-10$. Superconvergent decay starts only once N crosses the band, whose position is set by the oscillation content $\omega_{\max} T / 2\pi$ of the system — not by the method order; on the hard system the single-step curve ($p = 1$) cannot reach the band within the practical stage range and stays meaningless throughout.

block, factorization $\sim d^3 s^3$) and loses its order advantage entirely for non-smooth coefficients (act-and-wait, interrupted cutting).

4. **Explicit steps only for cheap low-accuracy sweeps on non-stiff systems.** Their observed order is capped by the available dense output (≈ 3 for classical RK4), and on stiff high-dimensional systems they are unusable below a CFL-like resolution bound, where the A-stable Gauss stages converge from the coarsest grids.
5. **Non-smooth coefficients: align or refine.** If switching instants can be placed on grid nodes, full order is retained; if not, expect order collapse and prefer moderate order with more steps.

Appendix D demonstrates the recipe end-to-end on two production stability charts, and Appendix E collects structural remarks on delay-period ratios, sub-step delays and mass matrices.

7 Conclusion

Once the stability analysis of time-periodic DDEs is written in multiplication-free banded form, the per-step integrator is a free design choice: Butcher-tableau residual blocks with continuous-extension delay interpolation raise the convergence order to that of the chosen scheme, up to Gauss superconvergence $2s$, for mesh-commensurate delays (orders between $s+1$ and $2s$ otherwise), while retaining $\mathcal{O}(p)$ complexity. The two classical method families are the endpoints of one continuum — piecewise-constant semi-discretization at maximal sparsity and global collocation at maximal per-step order — and neither endpoint is optimal: the CPU-time optimum sits at a problem-dependent sweet spot of moderate order and many steps, whenever the resolution demand or the system dimension is nontrivial; small smooth problems favour the near-spectral corner. The optimum drifts only logarithmically with the accuracy target and towards lower order for larger system dimension and non-smooth coefficients. This unifies and clarifies, rather

than reinvents; the actionable outputs are the sweet-spot guidance and the open-source SODS.jl implementation.

Declaration of conflicting interests

The author declared no potential conflicts of interest with respect to the research, authorship, and/or publication of this article.

Funding

This work was supported by the Hungarian Scientific Research Fund (grant number NKFIH-AG-152125).

A Model parameters

Delayed Mathieu equation ($d = 2$)

$$\ddot{x} + a_1 \dot{x} + (\delta + \epsilon_M \cos(2\pi t/T))x = b_0 x(t - \tau) + \sin(4\pi t/T),$$

with $\delta = 3$, $\epsilon_M = 0.2$, $b_0 = -0.15$, $a_1 = 0.1$, $\tau = T = 2\pi$. Converged dominant multiplier $\mu = 0.351417260953549$ (two-resolution GL10 reference, Appendix F).

Seasonal scalar model ($d = 1$) $\dot{y}(t) = -(a + b \cos t)y(t - \tau)$ with $a = b = 1$, $\tau = 2$, $T = 2\pi$; $\mu = 0.626021929290183$.

Turning with spindle-speed variation ($d = 2$)

$$\ddot{x} + \zeta \dot{x} + (1 + k_w)x = k_w x(t - \tau(t)) + \cos(2\pi t/T),$$

$$\tau(t) = \frac{2\pi}{\Omega} \left(1 + A_{SSV} \sin \frac{2\pi t}{T} \right),$$

with $k_w = 0.2$, $\zeta = 0.1$, $\Omega = 0.3$, $A_{SSV} = 0.1$, $T = N_T 2\pi/\Omega$, $N_T = 10$ (so $T/\tau_{\max} \approx 9.1$); $\mu = 4.476168383091$.

FEM beam with delayed boundary feedback ($d = 28$) Longitudinal rod of length $L = 100$ m, Young's modulus $E = 210$ GPa, cross-section $A = 10^{-4}$ m², density $\rho_m = 7800$ kg/m³, stiffness-proportional damping $\eta = 0.01$, discretized into $N_e = 15$ lumped-mass points with the first point fixed, leaving 14 free nodes ($d = 2 \cdot 14 = 28$ states; the predecessor's appendix quotes $d = 30$ for a slightly different

node counting of the same model). Delayed displacement feedback from the first free node to the last node with gain $PE/(L/N_e)$, $P = 0.2$; wave speed $c_0 = \sqrt{E/\rho_m}$, $T_{sp} = L/c_0$, delay $\tau = 0.2 T_{sp}$, period $T = 0.4 T_{sp}$ ($\tau/T = 1/2$); constant unit force at the free end. In the act-and-wait configuration the feedback matrix is active only for $t < 0.8 T$. Reference multipliers: $\mu = 373.008322$ (smooth) and $\mu = 319.636074$ (act-and-wait, aligned-grid reference).

B Additional convergence and work–precision studies

Figure 7 isolates the delay-commensurability effect quoted in Section 5.2. For the two systems whose delayed arguments never coincide with grid nodes — the seasonal model, where $\tau/\Delta t = p/\pi$ is never an integer, and the SSV turning system with its time-varying delay — the relative error of the dominant multiplier is re-plotted against p for the Gauss stages alone. Each method carries a shaded guide band spanning its interpolation-limited guaranteed order $s+1$ and its superconvergent cap $2s$, anchored at the finest pre-floor data point. Every measured curve stays within its band. On the seasonal model the fitted slopes (2.9 for GL2, 3.8 for GL3) press against the $s+1$ edge — the continuous-extension cap is active — while on the SSV system partial error cancellation keeps the observed orders higher: GL2 retains the full $4.0 = 2s$, GL3 attains 5.3, and GL5 attains 7.2, both strictly inside $(s+1, 2s)$. The one-stage GL1, whose bracket is degenerate ($s+1 = 2s = 2$), converges at second order regardless of commensurability. The GL5 tail flattens at $\varepsilon \approx 10^{-13}$, the agreement level of the two-resolution reference for this system; the plateau is displayed but excluded from the slope fit. Together with the mesh-commensurate orders of Table 1, these measurements substantiate the bracket $[s+1, 2s]$ stated in Section 5.2: $s+1$ is the guaranteed floor, $2s$ the attainable ceiling, and where a particular problem lands between them is a property of the problem, not of the method.

Figures 9–11 repeat the three-panel work–precision study of Figure 3 for the remaining test systems. The qualitative picture is identical: linear time complexity for every embedded integrator, Gauss stages dominating beyond $\varepsilon \approx 10^{-4}$. On the stiff beam the explicit methods are unstable below a CFL-like resolution bound (errors up to 10^{43}), while the A-stable Gauss stages converge from the coarsest grids. Figure 8 shows the corresponding (s, p) -maps for the low-demand Mathieu system.

C Non-smooth coefficients: act-and-wait beam

With act-and-wait switching, the beam’s feedback matrix jumps at $t = 0.8 T$. On grids where the switching instant is a node (p divisible by 5) the nominal order is retained (measured slopes 2.0 for GL1 and 3.7 for GL2; GL3 sits at the reference floor from the coarsest aligned grid). On off-grid meshes (powers of two) the observed order of *every* method collapses to ≈ 1 (measured 1.2, 1.1, 0.9 for GL1–GL3; Figure 12), confirming the scope remark of Section 5.7: with unaligned non-smoothness, all methods fall back to first

order, so nothing is lost by choosing the sparse moderate-order variants.

D Brute-force stability charts at the optimal setting

As an end-to-end demonstration of the sweet-spot guidance, high-resolution stability charts are computed by brute force — one eigenvalue solve per grid point (each solve single-threaded), using the CPU-optimal setting of each system. Figure 13 shows the results, coloured by the logarithm of the spectral radius (blue: stable, red: unstable); the stability boundary $\rho = 1$ is refined to sub-cell accuracy by the multi-dimensional bisection method (MDBM, Bachrathy and Stépán 2012) on top of the colour map. The grid evaluation is trivially parallel (every point is an independent solution-operator solve) and is executed on 16 threads. The delayed Mathieu chart ($100 \times 100 = 10000$ points, GL7, $p = 2$) completes in 3 s wall time plus 2 s for the MDBM boundary; the strongly modulated SSV turning chart ($100 \times 100 = 10000$ points, GL4, step count chosen per point by the sampling recipe — a resolution-pinned hard system) in 55 s wall time plus 25 s for the boundary. No continuation tricks are used; the per-point cost is simply that of the optimally chosen solution-operator solve.

E Structural remarks: delay–period ratios, sub-step delays, and mass matrices

Period-to-delay ratio. The ratio $T/\tau_{\max}$ (i.e. p versus r) never affects the triangularity of Φ_L ; it only selects the splitting into the left/right pair (the $p - 1 \geq r$ partitioning of Bachrathy 2026), and for $\tau_{\max} > T$ the assembly simply wraps the overlapping part of the history through unchanged — still triangular. The turning example ($T/\tau_{\max} \approx 9.1$) exercises the $p \gg r$ case throughout this paper.

Sub-step delays. What *does* change the structure is the delay relative to the *step*: for $\tau_k(t) \geq \Delta t$ every delayed lookup references strictly earlier blocks and Φ_L is unit block-lower-triangular. For $\Delta t > \tau_k$ stage lookups reach into the *current* step, so the diagonal block becomes a non-identity but small and invertible matrix: block forward substitution still applies at linear cost, although even an explicit tableau then requires a per-step solve (the classical “overlapping” phenomenon of DDE integration). Entries *above* the diagonal would require advanced-type arguments ($\tau < 0$) and never arise for retarded systems.

Mass matrices and differential–algebraic systems. The descriptor form $M\dot{\mathbf{x}} = \mathbf{A}(t)\mathbf{x} + \sum_k \mathbf{B}_k\mathbf{x}(t - \tau_k) + \mathbf{c}(t)$ replaces the identity blocks of the stage system (7) by $\mathbf{M}$:

$$\left[\mathbf{I}_s \otimes \mathbf{M} - \Delta t (\mathbf{a} \otimes \mathbf{I}_d) \mathcal{A}_n \right] \mathbf{Y} = (\mathbf{1}_s \otimes \mathbf{M}) \mathbf{x}_n + \dots \quad (11)$$

For regular $\mathbf{M}$ this is a drop-in replacement (the per-step LU absorbs $\mathbf{M}$; no explicit inverse is ever formed). For *singular* $\mathbf{M}$ — a delay differential–algebraic system (delay DAE) — the update row would require $\mathbf{M}^{-1}$, unless the tableau is *stiffly accurate* ($c_s = 1$, $b_j = a_{sj}$: Radau IIA, Lobatto IIIA), in which case $\mathbf{x}_{n+1} = \mathbf{Y}_s$ and no inversion occurs; the generalized eigenvalue




Figure 7. Convergence of the dominant multiplier for non-mesh-commensurate delays: seasonal model (left) and SSV turning system (right). Shaded bands span the order bracket $[s+1, 2s]$ per method (dotted edges), anchored at the finest pre-floor point; legend entries give the fitted log-log slope. All curves remain inside their bands: the seasonal model presses the interpolation-limited $s+1$ edge, whereas the SSV system retains orders up to the superconvergent cap. The GL5 plateau at $\varepsilon \approx 10^{-13}$ marks the accuracy of the two-resolution reference, not of the method.




Figure 8. The (s, p) -map for the low-demand delayed Mathieu equation (same three panels as Figure 5): here deep accuracy is reached in the high-order/low- p region as well, and the optima sit near the spectral corner because the matrices stay tiny; the polynomial CPU growth along s is nevertheless visible at the extended stage orders.




Figure 9. Work-precision study for the seasonal scalar model; same panel logic as Figure 3, parameters in Appendix A.

formulation also maps the algebraic part to multipliers near zero naturally. As a minimal validation (implemented in `examples/dae.mass.matrix.jl`), the index-1 system

$$\dot{x}_1 = -x_1 + x_2 + \frac{1}{2}x_1(t-1), \quad \nu \dot{x}_2 = \sin(t)x_1 - x_2,$$

with $\mathbf{M} = \text{diag}(1, \nu)$, reduces for $\nu = 0$ to a scalar linear DDE whose multiplier serves as reference. Radau IIA ($s = 3$) applied to the *singular* problem converges to this reference

at observed order approaching $4 = s+1$ (fitted slope 3.7 over the measured range), consistent with the delay-interpolation cap of Section 5.2 on this non-commensurate delay. The relative errors are $3.8 \cdot 10^{-8}$, $4.0 \cdot 10^{-9}$, $3.4 \cdot 10^{-10}$ and $1.6 \cdot 10^{-11}$ for $p = 25, 50, 100, 200$, and the singularly perturbed case $\nu = 10^{-6}$ converges stably to the DAE limit within the expected $\mathcal{O}(\nu)$ offset ($\approx 10^{-6}$). A full




Figure 10. Work–precision study for the turning system with spindle-speed variation; same panel logic as Figure 3, parameters in Appendix A.




Figure 11. Work–precision study for the FEM beam with delayed boundary feedback. The error axis is clipped at $3 \cdot 10^0$: below the CFL-like stability bound the explicit methods produce meaningless values (up to 10^{43}), which carry no information beyond “unstable”. CPU times are minima of repeated warm-started benchmark measurements.




Figure 12. Act-and-wait beam: relative error versus p for GL1–GL3 on grid-aligned (solid) and off-grid (dashed) meshes.

convergence theory for delay DAEs is beyond the scope of this paper.

F Reproducibility

All computations were performed with Julia 1.12.4 on Windows 11 (AMD Ryzen 7 5800H, 8 cores/16 threads, 16 GB RAM), with BLAS restricted to a single thread, package versions pinned by the repository manifests (SOSD.jl v0.1.1 with KrylovKit, LinearMaps, RungeKutta, StaticArrays; baseline SemiDiscretizationMethod.jl v0.4.2). Timings are

minima over repeated warm-started runs (repetition counts and scatter are recorded in the results files); start vectors of all Krylov iterations are deterministic. Reference multipliers are accepted only when two resolutions of a GL10 (small systems) or matrix-free GL5 (beam) computation agree to 10^{-10} relative, and the reference method is cross-validated against `SemiDiscretizationMethod.jl` on every system. The 10^{-12} -target optima lie at the limit of this verification and are indicative. Every figure can be regenerated via `julia --project=benchmark benchmark/run-all.jl` followed by `julia --project=benchmark benchmark/make-figures.jl`. Repository: <https://github.com/bachrathyd/SOSD.jl> [insert commit hash at submission].

References

- Altintas Y (2012) *Manufacturing Automation: Metal Cutting Mechanics, Machine Tool Vibrations, and CNC Design*. Cambridge University Press.
- Bachrathy D (2026) Linear time complexity analysis of time-periodic delayed systems with multiplication free semi-discretization method. *Journal of Vibration and Control* In press.
- Bachrathy D and Stépán G (2012) Bisection method in higher dimensions and the efficiency number. *Periodica Polytechnica Mechanical Engineering* 56(2): 81–86. DOI:10.3311/pp.me.




Figure 13. Brute-force stability charts at the optimal settings, coloured by $\log_{10} \rho$ (blue: stable, red: unstable); black dots: the $\rho = 1$ boundary refined by MDBM; the measured CPU times appear in each title. Left: delayed Mathieu equation in the (δ, b_0) plane ($\epsilon_M = 1$, $a_1 = 0.1$, $\tau = T = 2\pi$), GL7 with $p = 2$. Right: SSV turning in the (Ω, k_w) plane with strong modulation ($\zeta = 0.02$, $A_{SSV} = 0.3$, $N_T = 10$); here the step count follows the sampling recipe of Section 6 per point, $p(\Omega) = \lceil 16/\Omega \rceil$ — the paper's own rule applied in a realistic production setting. Grid evaluation on 16 threads; the MDBM refinement is serial.

2012-2.01.

- Bellen A and Zennaro M (2003) *Numerical Methods for Delay Differential Equations*. Oxford: Oxford University Press. DOI: 10.1093/acprof:oso/9780198506546.001.0001.
- Breda D, Maset S and Vermiglio R (2005) Pseudospectral differencing methods for characteristic roots of delay differential equations. *SIAM Journal on Scientific Computing* 27(2): 482–495. DOI:10.1137/030601600.
- Breda D, Maset S and Vermiglio R (2015) *Stability of Linear Delay Differential Equations: A Numerical Approach with MATLAB*. New York: Springer. DOI:10.1007/978-1-4939-2107-2.
- Butcher JC (2016) *Numerical Methods for Ordinary Differential Equations*. 3rd edition. Chichester: Wiley. DOI:10.1002/9781119121534.
- Haegeman J (2021) Krylovkit.jl: A julia package for krylov methods. GitHub.
- Hairer E and Wanner G (1996) *Solving Ordinary Differential Equations II: Stiff and Differential-Algebraic Problems*. 2nd edition. Berlin: Springer. DOI:10.1007/978-3-642-05221-7.
- Inspurger T (2006) Act-and-wait concept for continuous-time control systems with feedback delay. *IEEE Transactions on Control Systems Technology* 14(5): 974–977. DOI:10.1109/TCST.2006.876938.
- Inspurger T and Stépán G (2002) Semi-discretization method for delayed systems. *International Journal for Numerical Methods in Engineering* 55(5): 503–518.
- Inspurger T and Stépán G (2011) *Semi-discretization for Time-delay Systems*. Springer.
- Inspurger T, Stépán G and Turi J (2008) On the higher-order semi-discretizations for periodic delayed systems. *Journal of Sound and Vibration* 313(1-2): 334–341. DOI:10.1016/j.jsv.2007.11.040.
- Long XH and Balachandran B (2007) Stability of milling process with random spindle speed variation. *International Journal of Machine Tools and Manufacture* 47(10): 1600–1609.
- Munoa J et al. (2016) Chatter suppression techniques in metal cutting. *CIRP Annals* 65(2): 785–808.
- Seguy S et al. (2008) High-speed multiple spindle drilling stability with random perturbations. *Journal of Sound and Vibration* 311(3-5): 833–849.
- Stépán G (1989) *Retarded Dynamical Systems: Stability and Characteristic Functions*. Harlow: Longman Scientific & Technical.
- Tobias SA (1965) *Machine-Tool Vibration*. Blackie.